



Data Science Project Guide: Netflix

TechAcademy e.V.

Wintersemester 2022/23

Contents

1	Welcome!	4
2	What's Data Science and How Do I Do It?	5
2.1	What's R?	5
2.1.1	RStudio	5
2.1.2	DataCamp courses and Curriculum	5
2.1.3	Helpful Links	6
2.2	What's Python?	6
2.2.1	Anaconda and Jupyter	7
2.2.2	DataCamp courses and Curriculum	7
2.2.3	Helpful Links	9
2.3	Your Data Science Project	9
2.3.1	Prelude	9
2.3.2	Coding Meetups	9
2.3.3	Final Deadline	10
3	Introduction to Your Project	11
3.1	Purpose of the Project Guide	11
3.2	What is this Project About?	11
3.3	Exploratory Data Analysis – Getting to Know the Data Set	11
3.4	Recommendation System	11
4	Exploratory Data Analysis	12
4.1	Getting started	13
4.1.1	Discovering the Data	13
4.1.2	Give some overall statements	13
4.2	Data Cleaning and Useful Transformations	14
4.2.1	Date Formatting	14
4.2.2	More details of the longest movie	14
4.2.3	A histogram of the top 10 longest movies duration	14
4.2.4	Visualizing movie durations over time	15
4.3	Your personal data	15

4.3.1	Load your data	15
4.3.2	Clean and transform dataset	15
4.3.3	Merging datasets: primitive approach	15
4.3.4	Merging datasets: advanced approach	16
4.3.5	Dynamic line plot	16
4.4	Let's get personal	16
4.4.1	Monthly viewing time in 2021	17
4.4.2	Average per weekday	17
4.5	Binge watching	17
4.6	Scatterplot with marginal density	18
4.7	Word cloud with your favorite genre	18
4.8	Surprise Us!	18
5	Exercise Checklist	19
5.1	Part 1: Exploratory Data Analysis (Beginner + Advanced Tracks)	19
5.2	Part 2: Recommendation System(motivated Beginner + Advanced Tracks)	19
6	What's Next in Your Data Science Career?	20
6.1	Continue your journey on DataCamp	20
6.2	Data Science in General	20
6.3	R	20
6.4	Python	21

1 Welcome!

Let us start with an overview of what is waiting for you in the next weeks and in this project guide! If you are a beginner, In the first few chapters you will be introduced to the basics of the **R** and **Python** tracks respectively and you will find helpful explanations to questions you might have in the beginning of your coding journey. There will be a quick introduction to the Data Science track so that you can get started with the project quickly. So let's get started with the basics! You will work on your project in small groups of fellow students. This not only helps you get the project done faster, it also helps make your results even better. Our experience shows: Contrasting university backgrounds and different opinions and ideas will produce the best results. Besides, it is of course more fun to work on a project together than to code alone! The groups can consist of a maximum of four members and you can choose your teammates independently. We explicitly encourage you to collaborate with students from different departments. When submitting your final project, it is important to note: for a certificate, each person must submit the project individually. However, this can be identical within your group. You can get more information at our "Introduction to Coding" Workshop on November 9, 2022 or at our first Coding Meetup on November 23, 2022. This Netflix case study and the associated project guide was developed and written entirely from scratch by TechAcademy's Data Science team. Jessica Weigel and Isabel Schnorr developed the project in **R**, while Lea Karoza, Moritz Schwerdt and Rocky Auer developed it in **Python**.

2 What's Data Science and How Do I Do It?

2.1 What's R?

R is a programming language that was developed by statisticians in the early 90s for use in the calculation and visualization of statistical applications. R is now one of the most widely used programming languages in the field of data science. Code in R does not have to be compiled, but can be used interactively and dynamically. This makes it possible to quickly gain basic knowledge about existing data and to display it graphically. R offers much more than just programming, but also a complete system for solving statistical problems. A large number of packages and interfaces are available, with which the functionality can be expanded and integration into other applications is made possible.

2.1.1 RStudio

Install R and RStudio Locally To get started you should install R and RStudio locally on your computer. Before we try to put in word how to do so: Here's a [DataCamp tutorial](#) on how to do that. If you struggle, we will help you with installing at the "Introduction to Coding" Workshop or at our first Coding Meetup Until then, focus on learning the "hard skills" of programming with the courses on DataCamp. That brings us to your curriculum in the next section!

2.1.2 DataCamp courses and Curriculum

The following list shows the required DataCamp courses for the Data Science with R Track at TechAcademy. As a beginner, please stick to the courses of the "beginner" program. Ambitious beginners can, of course, take the advanced courses afterward. However, it would be best if you worked through the courses in the order we listed them.

The same applies to the advanced courses. Here, too, you should finish the specified courses in the given order. Since it can, of course, happen that you have already mastered the topics of an advanced course, you can replace some courses. If you are convinced that the course does not add value to you, feel free to replace it with one of the courses in the "Exchange Pool" (see list below). However, you should not pursue an exchange course until you finish all chapters from the advanced course: "Intermediate R."

To receive the certificate, both beginners and advanced learners must complete at least 6 courses of the curriculum (6/9 courses). For the beginners, this means until – and including – the course [Cleaning Data in R \(4h\)](#) and for the advanced until –and including – [Modeling with Data in the Tidyverse \(4h\)](#). After completing the curriculum and the project's requirements, you will receive your TechAcademy certificate!

**Data Science in R Fundamentals (Beginner)**

1. [Introduction to R \(4h\)](#)
2. [Intermediate R \(6h\)](#)
3. [Data Manipulation with dplyr (4h)] (<https://www.datacamp.com/courses/data-manipulation-with-dplyr-in-r>)
4. [Cleaning Data in R (4h)] (<https://www.datacamp.com/courses/cleaning-data-in-r>)
5. [Introduction to Data Visualization with ggplot2 (4h)] (<https://www.datacamp.com/courses/data-visualization-with-ggplot2-1>)
6. [Exploratory Data Analysis in R (4h)] (<https://www.datacamp.com/courses/exploratory-data-analysis>)
7. [Reporting with R Markdown (4h)] (<https://app.datacamp.com/learn/courses/reporting-with-rmarkdown>)

Data Science in R (Advanced)

1. [Intermediate R \(6h\)](#)
2. [Data Manipulation with dplyr (4h)] (<https://www.datacamp.com/courses/data-manipulation-with-dplyr-in-r>)
3. [Cleaning Data in R (4h)] (<https://www.datacamp.com/courses/cleaning-data-in-r>)
4. [Introduction to Data Visualization with ggplot2 (4h)] (<https://www.datacamp.com/courses/data-visualization-with-ggplot2-1>)
5. [Text Mining with Bag-of-Words in R (4h)] (<https://app.datacamp.com/learn/courses/text-mining-with-bag-of-words-in-r>)
6. [Writing Efficient R Code \(4h\)](#)
7. [Reporting with R Markdown (4h)] (<https://app.datacamp.com/learn/courses/reporting-with-rmarkdown>)

Data Science in R (Advanced) – Exchange Pool

- [Intermediate Data Visualization with ggplot2](#)
- [Interactive Maps with leaflet in R \(4h\)](#)
- [Interactive Data Visualization with plotly in R \(4h\)](#)
- [Multiple and Logistic Regression in R \(4h\)](#)
- [Machine Learning in Tidyverse \(5h\)](#)
- [Introduction to Writing Functions in R \(4h\)](#)

2.1.3 Helpful Links

- [RStudio Cheat Sheets](#)
- [RMarkdown Explanation](#) (to document your analyses)
- [StackOverflow](#) (forum for all kinds of coding questions)
- [CrossValidated](#) (Statistics and Data Science forum)

2.2 What's Python?

Python is a dynamic programming language. You can execute the code in the interpreter, so you do not have to compile the code first. This feature makes Python very easy and quick to use. The excellent usability, easy readability, and simple structuring were and still are core ideas in developing this programming language.

You can use `Python` to program according to any paradigm, whereby structured and object-oriented programming is most straightforward due to the structure of the language. Still, functional or aspect-oriented programming is also possible. These options give users significant freedom to design projects the way they want and great space to write code that is difficult to understand and confusing. For this reason, programmers developed specific standards based on the so-called `Python` Enhancement Proposals (PEP) over the decades.

2.2.1 Anaconda and Jupyter

Before you can use `Python`, you must install it on the computer. `Python` is already installed on Linux and Unix systems (such as macOS), but often it is an older version. Since there are differences in the handling of `Python` version 2 – which is no longer supported – and version 3, we decided to work with version 3.6 or higher.

One of the easiest ways to get `Python` and most of the best-known programming libraries is to install Anaconda. There are detailed explanations for installing all operating systems on the [website](#) of the provider.

With Anaconda installed, all you have to do is open the Anaconda Navigator, and you're ready to go. There are two ways to get started: Spyder or Jupyter. Spyder is the integrated development environment (IDE) for `Python` and offers all possibilities from syntax highlighting to debugging (links to tutorials below).

The other option is to use Jupyter or Jupyter notebooks. It is an internet technology-based interface for executing commands. The significant advantage of this is that you can quickly write shortcode pieces and try them out interactively without writing an entire executable program. Now you can get started!

If you have not worked with Jupyter before, we recommend that you complete [this DataCamp course](#) first. There you will get to know many tips and tricks that will make your workflow with Jupyter much easier.

To make your work and, above all, the collaboration more accessible, we are working with the [Google Colab](#) platform that contains a Jupyter environment with the necessary libraries. You can then import all the data required for the project with Google Drive. We will introduce this environment during our first Coding Meetup. Until then, focus on learning the “hard skills” of programming with your courses on DataCamp. This topic brings us to your curriculum in the next section!

2.2.2 DataCamp courses and Curriculum

The following list shows the required DataCamp courses for the Data Science with `Python` Track at TechAcademy. As a beginner, please stick to the courses of the “beginner” program. Ambitious beginners can, of course, take the advanced courses afterward. However, it would be best if you worked through the courses in the order we listed them.

The same applies to the advanced courses. Here, too, you should finish the specified courses in the given order. Since it can, of course, happen that you have already mastered the topics of an advanced course, you can replace some courses. If you are convinced that the course does not add value to you, feel free to replace it with one of the courses in the “Exchange Pool” (see list below). However, you should not pursue an exchange course until you finish all chapters from the advanced course: “Intermediate Python.”

To receive the certificate, both beginners and advanced learners must complete at least two-thirds of the curriculum (6/9 courses). For the beginners, this means until – and including – the course “[Joining Data with pandas \(4h\)](#)” and for the advanced until –and including – “[Introduction to Linear Modeling in Python \(4h\)](#).” In addition, you should complete at least *two-thirds* of the project tasks. After completing the curriculum and the project’s (minimal) requirements, you will receive your TechAcademy certificate!



Data Science with Python Fundamentals (Beginner)

1. [Introduction to Data Science in Python \(4h\)](#)
2. [Intermediate Python \(4h\)](#)
3. [Python for Data Science Toolbox \(Part 1\) \(3h\)](#)
4. [Introduction to Data Visualization with Matplotlib \(4h\)](#)
5. [Data Manipulation with pandas \(4h\)](#)
6. [Joining Data with pandas \(4h\)](#)
7. [Exploratory Data Analysis in Python \(4h\)](#)
8. [Working with Dates and Times in Python \(4h\)](#)
9. [Introduction to Importing Data in Python \(3h\)](#)

Data Science with Python (Advanced)

1. [Intermediate Python \(4h\)](#)
2. [Python Data Science Toolbox \(Part 1\) \(3h\)](#)
3. [Python Data Science Toolbox \(Part 2\) \(4h\)](#)
4. [Cleaning Data in Python \(4h\)](#)
5. [Exploratory Data Analysis in Python \(4h\)](#)
6. [Introduction to Linear Modeling in Python \(4h\)](#)
7. [Statistical Thinking in Python \(Part 1\) \(3h\)](#)
8. [Time Series Analysis in Python \(4h\)](#)
9. [Machine Learning for Time Series Data in Python \(4h\)](#)

Data Science with Python (Advanced) - Exchange Pool

- [Interactive Data Visualization with Bokeh \(4h\)](#)
- [Data Visualization with Seaborn \(4h\)](#)
- [Supervised Learning with scikit-learn \(4h\)](#)
- [Linear Classifiers in Python \(4h\)](#)
- [Unsupervised Learning in Python \(4h\)](#)
- [Introduction to Deep Learning in Python \(4h\)](#)
- [ARIMA Models in Python \(4h\)](#)
- [Web Scraping in Python \(4h\)](#)
- [Writing Efficient Python Code \(4h\)](#)
- [Writing Efficient Code with pandas \(4h\)](#)

2.2.3 Helpful Links

Official Tutorials/Documentation:

- <https://docs.python.org/3/tutorial/index.html>
- <https://jupyter.org/documentation>

Further Explanations:

- <https://pythonprogramming.net/>
- <https://automatetheboringstuff.com/>
- <https://www.reddit.com/r/learnpython>
- <https://www.datacamp.com/community/tutorials/tutorial-jupyter-notebook>

2.3 Your Data Science Project

Participant's Section:

- <https://www.tech-academy.io/teilnehmerbereich>

2.3.1 Prelude

Now that you have learned the theoretical foundation of Data Science in the DataCamp courses, you can put your skills into practice. We have put together a project for you based on real data sets. You can read about the details of this project in the following chapters of this project guide. If you still feel a little unsure about how to start off, there will be a Coding Introduction on November 9, 2022, where we will give you a general sense on how to start off with Python or R. Of course, we will also describe the project. We will discuss everything you need to know during the first Coding Meetup, which will take place on November 23, 2022. After that, your work on the project will officially begin. You can find the exact project tasks together with further explanations and hints in the following chapters.

2.3.2 Coding Meetups

To give you a little overview on the dates of our Meetups, you can find all the dates in one go here! - **02.11.2022**, Kick-Off Event - **09.11.2022**, Introduction to Coding Workshop - **23.11.2022**, Coding MeetUp 1 - **14.12.2022**, Coding MeetUp 2 - **11.01.2023**, Coding MeetUp 3 - **05.02.2023**, Coding MeetUp 4

To receive the certificate, it is essential that you hand in your project 4 times during the semester. The following dates are due dates:

2.3.3 Final Deadline

For both your DataCamp courses and the project files hand-in the deadline is **05.02.2023, 23:59**

3 Introduction to Your Project

3.1 Purpose of the Project Guide

This document will guide you through the different steps of your project and will provide you with valuable hints along the way. However, it is not a detailed step-by-step manual because we felt like you needed to develop the skills of coming up with your way of solving different tasks. This method is a great way to apply the knowledge and tools you have acquired through DataCamp. Since data science concepts are independent of specific programming languages, we will describe the general approach in a text chunk. Having understood the bigger picture and starting with the tasks, you will find language-specific tips and tricks in visually separated boxes (R -Track : blue-bordered boxes, Python - Track : yellow-bordered boxes). Questions might come up, or you might not know how to solve a task right away—but don't worry—this is just part of coding. In those cases, you could ask your team members for help. If they are not able to help, and in the unlikely case that even Google can't help you, the TechAcademy mentors will help you via Slack or directly during the coding meetups. At the end of the project guide, you will find an overview of all tasks that have to be completed, depending on your track (beginner/advanced). You can use this list to check which tasks you still need to complete or which assignments are relevant for your track.

3.2 What is this Project About?

This semester, we will have a look at Netflix data! More precisely, we are first analyzing a very detailed general Netflix data set, compiled on [kaggle](#), after which you will get the chance to look at your own Netflix data. You will find all kinds of information in the data sets - valuable and useless ones. Are you already curious to see for yourself? In analogy to the typical Data Science workflow, we split this project into two parts.

3.3 Exploratory Data Analysis – Getting to Know the Data Set

This first part of the project is structured in a way that lets you get to know the data thoroughly by completing the given tasks one after the other. As a beginner, you can stop after this part because you will have fulfilled the necessary coding requirements for the certificate. However, if this first part inspires you to learn more, we encourage you to also work on the second part. If you get stuck, **Google** and [StackOverflow](#) are amazing problem solver (besides the mentors, of course).

3.4 Recommendation System

This part is mainly for the advanced TechAcademy participants. If you are a beginner and you were able to complete the first part without too many difficulties, we highly recommend trying to do the second part as well.

4 Exploratory Data Analysis

Before you can dive into the data, you need to do two things. First of all, you need to request your personal data from Netflix, which can take a few days. Therefore, you should ask for your data as early as possible. (If you do not have access to a Netflix account, contact your mentor, and we will provide you with some data)

Download your data

1. Follow this link to request your data <https://www.netflix.com/account/getmyinfo>, and sign in with your username and password.
2. After that, click on the red box Submit Request.
3. Now you need to verify your request via Mail.
4. Once you confirm the request, you will be forwarded to the website and have to enter your password again. If your request is in progress you might have to wait a few days until the data is available to download.
5. When you receive the mail “Your download is ready,” you can download your data by clicking on the red box Get Information. Don’t wait too long since you only have 7 days to download it from your account!
6. After confirming your password, you can download your data as a netflix-report.zip file. Unzip the file, open the folder CONTENT_INTERACTION and look for the ViewingActivity.csv file.

However, you can already start with the first exercises of this chapter with data we provide. To start right away, you need to set up your programming environment. This will be the place where the magic happens.



We work with R and RStudio locally which requires a setup. If you haven’t installed it yet here is a [DataCamp tutorial](#). If you struggle, we will help you with installing at the “Introduction to Coding” Workshop or at our first Coding Meetup



We recommend using [Google Colab](#) for this project since it requires no particular setup, stores its notebooks to your Google Drive, and makes it easy for you to share them with your team members.

As an alternative to Google Colab, you might want to install Jupyter Notebook locally using the Anaconda distribution. We will give you a more detailed step-by-step demo during the “Introduction to Coding Workshop” or our first Coding Meetup.

4.1 Getting started

4.1.1 Discovering the Data

The first big step is importing a general Netflix data set into your coding environment: *

```
[Netflix]
```

Once done, let's start by looking at the Netflix data set.

Have a look at the columns of the data set and their "values" Do you see any missing values or data entries that are different from the other entries? *What are the data formats (e.g. data types)?* Look at the `date_added` columns. Do you see any critical aspects of the data? Write down a couple of sentences to these questions, the goal is to show (your readers) what you are seeing. Also, comment on any errors or irregularities which you notice and that could be an issue later on in the project.



You can import data files by, for examples, using the command `netflix_general <- read_csv("Netflix.csv")`. However, it has to be the correct path. Use `head()` and `glimpse()` or `class()`(on a specific column) to get an overview. Lastly, get a quick summary of the data using `{base} R`'s `summary()` or `{psych}`'s `describe()` function. To find more ways to generate quick data summaries in R, check this blog post from Adam Medcalf. If you need some additional, more general information on how to import data and different data types, check out this [cheat sheet](#).



You can feed the links to the respective data files above to a method of the [pandas package](#) (you might want to specify the index column). Check out the resulting `pd.DataFrame` instance with the `head()` method and the `dtypes` attribute. You can also dig into a specific column with `describe()`. Here's an additional [pandas cheat sheet](#) for you to reference

4.1.2 Give some overall statements

Now that you have imported the data set, let's have a deeper look. Since you already got a feeling by now, it would be interesting to indicate some outstanding features (based on the uncleaned and untransformed dataset)!

What's the longest movie (not TV show) included in the dataset? (Look very closely if your result is valid) Which country released the most content (movies and tv shows)? *How many movies and tv shows are included? Try to do a simple plot here so you can see the distribution. If you cannot plot it just now, come back later to finish this.



You could use `{dplyr}`'s `select`, `filter`, `count` and `arrange` function to compute the desired outcome to answer the question. If you haven't heard of the `{dplyr}` package yet, take the respective DataCamp course asap! For the plot, look at `{ggplot}` and then call `geom_bar()`. `{ggplot2}`'s syntax is perhaps a bit harder to learn at the start, but it gives you more plotting benefits in the long run. If you haven't heard of the `{ggplot2}` package yet, take the respective DataCamp course.

4.2 Data Cleaning and Useful Transformations

4.2.1 Date Formatting

From exploring the data in the previous tasks you might have noticed that the time and dates are not in a “date format”. In order to fix this, convert the “date_added” and “release_year” column into date format.



You can use the lubridate package to transform the column into a date format. To make things easier, check out the [lubridate](#) cheat sheet. Converting the release year might take two steps.



pd.to_datetime() is what I would look at for example. A corresponding DataCamp resource is section 4 in [Working with dates and times in Python](#). Also, the [Data Manipulation with pandas](#) course is of great help for the following exercises

4.2.2 More details of the longest movie

So far, you have converted the column into a more appropriate format for further investigations. Can you also fix the duration of the movies into a numeric format? Can you tell now what the longest movie is? On the concept of distribution, you should also try to compute the mean and standard deviation of the movie duration in minutes.



First off, you might want to create a separate data frame for the movies so the tv shows do not bother you further. Use the filter call here. Next you might have recognized the “min” appendix. This can be deleted with the “gsub” call. With the help of mutate you can convert the column into numeric values and arrange the movies descending. For the mean and standard deviation simply use {base} R’s mean() function to compute a mean of any column in R. Similarly, R comes with the sd() function to calculate the standard deviation of a column. You can combine these two functions, e.g., in a {dplyr}’s summarise() verb.

4.2.3 A histogram of the top 10 longest movies duration

Visualizing data is essential to facilitate perception and to understand information: Create a graph to visualize the top 10 longest movies. We focus on creating a histogram in the tips section since it’s perhaps the most common approach for plotting. You can, however, choose a different chart type; just make sure that the information you want to display is clear and correct. However, as in school, always add axis labels when possible.



If your data is saved as a dataframe you can perceive with the data from the previous task. Use {ggplot2} to create a histogram. Integrate top_n() before the ggplot() call and select geom_col() to achieve a histogram presentation.

4.2.4 Visualizing movie durations over time

While we are at the topic of movie durations, it might also be interesting to see how the average movie length evolved: Choose and plot the chart you think is most appropriate for displaying this type of information. After plotting, please comment and interpret the graph: Were there any significant increases/decreases in movie length over time? If so, what could be the reason?



Again, it's ggplot's turn. Instead of preparing the summarise() and mean() functions combination you can also directly integrate the adjustments in the geom_bar call with (stat = "summary", fun = "mean").

4.3 Your personal data

4.3.1 Load your data

You will now use the data set which you requested from Netflix. In the Netflix folder you will find the document of interest: ViewingActivity. Load this in your environment and inspect it as you did before with the Netflix Dataset. If you can't request your data, ask your mentor; they will provide you with an alternative data set.

4.3.2 Clean and transform dataset

As you might have noticed, Netflix recorded every time you clicked on a movie even if you didn't watch it. Check which column indicates those with a specific value. To avoid a bias in the following analysis, delete the respective rows. You can further drop columns that seem unnecessary or don't give any information and change column names if you wish. Additionally, convert the column Duration in minutes (round them) and extract the viewing day and viewing time from "Start time" into two separate columns (e.g., "date" and "time").



mutate and filter are your best friends in this cleaning process.

4.3.3 Merging datasets: primitive approach

As a data scientist, you'll often find yourself working with data sets from different data sources referencing the same object. For example, you might have the movie names in one file and the respective genre in a separate file. It would make more sense to just merge the two data sets into one. Indeed, this is the case with our data. Your Netflix data does not provide information on genre, actor, or director, while the general Netflix data set does. So, to make life easier for the upcoming tasks, you'll now need to merge both data sets by the title name.

There are several things kind of "wrong" with the merged data set. What is it and why? Tipp: You need to prepare the dataset in a way that the title, session and episode will be split into

a column of title, session and episode separately in order for the two data frames to join each other properly.



Calls such as `separate()`, `ifelse()` and `rename()` are suitable here. Merging can be done in several ways. You could, e.g., use `{dplyr}`'s `left_join()` to combine the two data sets based on the “left-sided” dataset. Which one should be the “left-sided”?



The merge function allows for different types of merges. Try coming up with the most logical one in this scenario!

4.3.4 Merging datasets: advanced approach

You might have noticed that merging the data set is a little challenging. This is because Netflix sometimes displays title, session and episode in a remixed way and also because your Netflix data might be in German whereas the general Netflix Data is in English. We fixed this issue externally and you can now send your data to Moritz Schwerdt, who will provide you with a respective general Netflix dataset. Now try merging again!

4.3.5 Dynamic line plot

Your goal for this task is to plot how each viewer's activity was recorded over time. Since it would be a bit unclear in a static plot, let's try to do a dynamic chart here!



The data is already ready for this. Your `ggplot()` call might need some more input than you used before. As the graphic is supposed to be a line with points `geom_line()` and `geom_point()` will be a good choice of combination. The dynamic aspect is no magic, you just need to add the `transition_reveal()` call.

4.4 Let's get personal

We dived into the provided general Netflix data set and a bit into your overall personal Netflix dataset, but now is the time to look into your very own Netflix history. First, we want to look at the longest movie you have ever watched, afterward, into your general and binge-watching behavior. We recommend using Tobi's data if you are currently working with the Netflix data set provided from us.

So, what's the longest movie you have ever watched? Maybe you also want to check the day? Do you remember that day?



`filter`, `select` and `arrange` will do a good job here.



Useful functions might be `sort_values()` and `groupby()`.

4.4.1 Monthly viewing time in 2021

Let's get a little fancy by displaying the monthly average time you watched in 2021 (or another year of your choice).



This task involves three steps. First, you want to extract the month from the “start_time” column, where `lubridate` will help you. You can then use the `group_by` argument to group the dataset by the month and filter for a specific year. After that, you can pipe along and directly add the third step: your `ggplot` call. Now, you can visualize your output however you want. The presented output is obtained by combining `geom_col` + `coord_polar()`.



We recommend a stepwise course of action: first, extract the month from your date column - at this point you should have a clue which package might be able to help! Then, you should groupby month and filter for a year of your choice. The last step is the generation of your plot! A great package for this task is [plotly.graph_objects](#). You might want to play around with the [Barpolar](#) function a little to receive the desired output.

4.4.2 Average per weekday

Our subsequent interest lies in analyzing the viewing time of specific weekdays. Either choose the whole year or a month of your interest and look at the weekday average. On which days have you watched more Netflix, can you see a peak? For our advanced programmers: Would it not be interesting to look at what time of the day you watch the most?



As before, you now want to extract the weekday from the “start_time”. Further you might want to filter for a specific month or year and group by date and weekday. You will then need to sum the minutes per day and group again by weekday to summarize for the mean of the duration. Finally, you could, for example, plot a classic `ggplot`.



The procedure is similar to before. Get the weekday from your data, compute the sum of daily minutes of watching, and generate a bar plot. The [seaborn package](#) is an amazing tool that generates beautiful plots with really limited code!

4.5 Binge watching

In this section, the goal should be to create a plot of your top 10 binge tv shows. Before you begin, you need to decide what binge watching means to you. We decided it's at least 1 hour of a TV show per day. But you should adjust that to your watching behavior.



Again, you could use the familiar `{dplyr}` verbs for this part. You first want to filter for TV Shows and group by date and title. You then need to count episodes and minutes per day and group again by title. Now, calculate the sum of episodes and minutes per day. Here you might want to convert minutes into hours. Finally, arrange everything descending. For the plot, you can use `ggplots geom_col()` or whatever you think is best to present your binge series

4.6 Scatterplot with marginal density

You have computed and visualized your favorite binge tv shows, but how has your watching behavior developed on Netflix since you first used it? This question is something we want to analyze now. There are different possibilities to explore this question. We will start by visualizing it via a scatterplot with marginal density. Include all the profile names for this task and make a visual comparison. What can the plot tell us about your watching behavior?



Again, seaborn is your best friend! The jointplot function will give you a great plot if you use it correctly.

4.7 Word cloud with your favorite genre

Until now, even though we merged the two data sets, analyzing the Netflix movies was possible by only using each one. Let's change that! Generate a word cloud with your personal most watched genres! Is it what you expected it to be?



Install the package "wordcloud" and start preparing your joined dataset. We are only interested in the "listed_in" column now. Use the mutate function to do a structure split, so each genre will get its row instead of being separated by a comma. Add unnest to your call. Extract the column and fix possible issues (e.g., white space). An easy way is to use the table() call on the listed_in and change the table directly into a data frame. For the word cloud, set a seed if you want to get the same cloud every time you run the code. The call wordcloud only needs to be filled with the respective word and frequency column. Feel free to spice it up!

4.8 Surprise Us!

This section only scratches the surface of what is possible with your Netflix data and some columns haven't even been analyzed. Have you discovered anything cool and/or want to experiment with a new visualization technique to deliver a message? Here is the right place to put down anything you did and would like to do with the project that did not quite fit in the other exercise sections. If you are looking for some visualization inspiration, check out this page. Congratulations! Based on your work with fundamental data transformations and many visualizations, you now have a solid understanding of the Netflix data sets and your personal data! With this, you have completed the EDA part of the project! Don't forget to send your project results to our project submission email address (projekt@tech-academy.io) before the deadline (05.02.2023, 23:59). Thanks for being a part of TechAcademy! If you are in the advanced track your coding journey goes on with the next section!

5 Exercise Checklist

This checklist should help you keep track of your exercises. Remember that you have to hand in satisfactory solutions to at least *two-thirds* of the exercises. If you're part of the beginner track, this refers to two-thirds of section 4 (EDA) only. If you're part of the advanced track, you have to hand in at least two-thirds of section 4 and 5. Hence, you'll need at least 66% in each of the two sections for a certificate.

5.1 Part 1: Exploratory Data Analysis (Beginner + Advanced Tracks)

You completed at least 6 courses on Datacamp.

You completed section 4 and handed your project at least 3 times in.

You checked that you included everything you coded in the right order and added explanatory comments.

You have all your output in the desired format.

You compared your output with your group.

You sent your final work to projekt@tech-academy.io (Remember: Each team member has to hand in their final work individually, in any case it can be the same within the group).

5.2 Part 2: Recommendation System (motivated Beginner + Advanced Tracks)

Everything of the above stated.

In addition you also completed section 5 and handed your project at least 3 times in.

6 What's Next in Your Data Science Career?

6.1 Continue your journey on DataCamp

You can continue learning on DataCamp until the start of the next Semester! Check out the [Skill](#) and [Career Tracks](#) and make sure to complete them to receive DataCamp's Statement of Accomplishment.

6.2 Data Science in General

Version Control with [Git](#)

Advice for [Non-Traditional Data Scientists](#)

Learn from Great Data Scientists on [Kaggle](#)

Great insights and helpful tips on [towardsdatascience](#)

6.3 R



Version Control with Git

RStudio has a nice interface that lets you enjoy the perks of Git without ever having to touch the command line – sounds great, does it? Learn how to set up the Git & R workflow with [Happy Git with R](#).

R Graph Gallery

Get inspiration to take your plotting to the next level. Includes code to reproduce the plots.

Follow the R Master Himself and the R Community

Hadley Wickham was and continues to be extremely influential on the development of R and its rise to one of the most popular data science languages. He's behind many tools that we taught you in this semester, especially the tidyverse (including great packages such as ggplot2 and dplyr). Follow him [on Twitter](#) to get great R advice and keep up to speed with everything new to R. Following the many people behind R (not only Hadley) is a great way for acquiring deeper understanding of the language and its developments.

Join the [Campus useR Group in Frankfurt](#)

There's a quite active R community in Frankfurt that meets once a month. It's open for students, professors, industry practitioners, journalists, and all people that love to use R. In those meetings, you'll hear about other's work, discuss new developments, and ask questions.

Listen to R Podcasts Another great way to easily keep up with new developments in the Data Science/R community. Check out [Not So Standard Deviations](#) or [the R-Podcast](#)

6.4 Python



Install Python Locally

Until now you've only programmed using JupyterHub on the TechAcademy Server. A next step would be to install Python and Jupyter locally on your computer. This [link](#) contains the necessary information on how to install the software on Windows, iOS or Linux.

Choosing the Right Editor

Using Jupyter is especially useful for short data analyses. But sometimes you want to write longer scripts in Python. In these cases, it is often more convenient to use a code editor instead of Jupyter. [This tutorial](#) highlights the positive aspects of such an editor and how to choose the right one for you. Pro-Tip: Also check out the other tutorials on [Real Python](#) and check out the Community Version of "PyCharm" which is the most common Python IDE (Integrated Development Environment).

Python Graph Gallery

Get inspiration to take your plotting to the next level. Includes code to reproduce the plots.

More Advanced Python Concepts

You know the basic data structures in Python like lists and dictionaries. What are the next steps to improve your knowledge? [This website](#) gives good explanations for slightly more advanced concepts which can be very useful from time to time.

A Deeper Understanding

If you want to get a deeper understanding of the Python programming language and into typical algorithms which are used in the field of Data Science, this [free book](#) can be a good starting point.

Writing Beautiful Python Code

"My code doesn't look nice, but it works!" This might work for yourself, but often you will work on code with other people. But even if you're just coding for yourself it's a good idea to follow the PEP8 style guide. It's a useful convention on how to structure and code in Python. You'll find useful resources for PEP8 [here](#) and [here](#).

Listen to Python Podcasts

When you don't have time for books you can listen to [Talk Python](#) or the [Python Podcast](#).